



DEVOPS CAPSTONE PROJECT

PROJECT 1 – INFRA OPTIMIZATION

WRITTEN BY: TAN KAH XUAN

LAST UPDATED: 2021-11-24

Contents

Introduction.....	2
Project 1 – Infra Optimization.....	2
Task 1 – Provisioning of EC2 instances with Ansible	3
Task 2 – Setting up a Kubernetes Cluster	6
Task 3 – Configuration of Kubernetes Clusters for Web Application	11
Setup Load Balancer	11
Setup MongoDB Database	11
Setup Web Application	15
Task 4 – Create Users in the Kubrnetes Cluster	20
Task 5 – Backup ETCD cluster data	21
Task 6 – Auto Scaling	22
Appendix – To reset kubernetes cluster.....	23

Introduction

Project 1 – Infra Optimization

Create a DevOps infrastructure for an e-commerce application to run on high-availability mode.

Background of the problem statement:

A popular payment application, EasyPay where users add money to their wallet accounts, faces an issue in its payment success rate. The timeout that occurs with the connectivity of the database has been the reason for the issue. While troubleshooting, it is found that the database server has several downtime instances at irregular intervals. This situation compels the company to create their own infrastructure that runs in high-availability mode.

Given that online shopping experiences continue to evolve as per customer expectations, the developers are driven to make their app more reliable, fast, and secure for improving the performance of the current system.

Implementation requirements:

1. Create the cluster (EC2 instances with load balancer and elastic IP in case of AWS)
2. Automate the provisioning of an EC2 instance using Ansible or Chef Puppet
3. Install Docker and Kubernetes on the cluster
4. Implement the network policies at the database pod to allow ingress traffic from the front-end application pod
5. Create a new user with permissions to create, list, get, update, and delete pods
6. Configure application on the pod
7. Take snapshot of ETCD database
8. Set criteria such that if the memory of CPU goes beyond 50%, environments automatically get scaled up and configured

Project Requirement Checklist

The Project will be splitted into 6 major tasks as listed below. The following tools will be involved in the setting up the project.

1. Ansible – The software provisioning tool for enabling infrastrucutre as code
2. AWS EC2 – The ubuntu environment nodes used to setup the Kubernetes Clusters
3. Docker – The container orchestration tool
4. Kubernetes – The containerized application management tool
5. Github – Source Code Repository

S/N	Tasks	Status
1	Provisioning 3 x AWS EC2 instance (with load balancer and elastic IP) with Ansible <ul style="list-style-type: none">• 1 x Kubernetes Master Node• 2 x Kubernetes Worker Node	Done
2	Install Docker and setup a Kubernetes Cluster	Done
3	Configuration of Kubernetes Cluster for Web Application <ul style="list-style-type: none">• Create Database Pod – MongoDB / Elasticsearch• Create Web Application Pod – NodeJS backend application• Implement relevant network polices at the database pod to allow ingress traffic from front-end application pod	Done
4	Configuration of the Pods <ul style="list-style-type: none">• Create a new user with permissions to create, list, get, update, and delete pods	Done
5	Taking a snapshot of ETCD database	Done
6	Auto Scaling – If the memory of CPU goes beyong 50%, environments automatically get scaled and configured	Done

Task 1 – Provisioning of EC2 instances with Ansible

Create new Ubuntu EC2 instance

The EC2 instance will be used to host ansible. There are several steps to prepare the EC2 instance.

1.a.1 Create AWS user

On the AWS console, search for IAM (Identity and Access Management) and create a new user. Here, take note of the name of the key pair generated which will be used by Ansible to setup the instances.

Create an IAM role with AmazonEC2FullAccess policy.

1.a.2 Provision an EC2 Instance

The t2.micro instance is launched. Attach the IAM Role with AmazonEC2FullAccess to the EC2 instance. Open port 22 for SSH.

Install ansible

1.b.1 Login to the EC2 instance using Putty/ MobaXterm or other tools to SSH into the ubuntu machine. Execute the following codes to install ansible on the ubuntu machine:

Install Ansible

```
$ sudo apt-get install -y ansible
```

```
$ sudo apt-get install python-pip -y
```

Install the python AWS SDK – Boto Framework

```
$ sudo pip install boto boto3
```

```
$ sudo apt-get install python-boto -y
```

Ansible will be using boto SDK to access various AWS resources.

Verify ansible installation

```
$ ansible --version
```

Create Ansible Playbook

1.c.1 Edit Ansible hosts or create an inventory file

```
$ sudo nano /etc/ansible/hosts
```

Add the contents at the end of the file:

```
[localhost]
localhost
```

1.c.2 Create Ansible playbook

```
$ mkdir playbook
```

```
$ cd playbook
```

```
$ sudo nano create_ec2.yml
```

```
---
- name: provisioning EC2 instances using Ansible
  hosts: localhost
  connection: local
  gather_facts: False
  tags: provisioning
```

```
vars:
```

```
keypair: 06Dec
instance type: t2.micro
image: ami-0907c2c44ea451f84
wait: yes
group: k8cluster
count: 3
region: ap-southeast-1
zones: ap-southeast-1b
security_group: k8cluster-security-group
```

tasks:

- name: Create security group
local_action:
 module: ec2_group
 name: "{{ security_group }}"
 description: Security Group for webserver Servers
 region: "{{ region }}"
 rules:
 - proto: tcp
 from_port: 22
 to_port: 22
 cidr_ip: 0.0.0.0/0
 - proto: tcp
 from_port: 8080
 to_port: 8080
 cidr_ip: 0.0.0.0/0
 - proto: tcp
 from_port: 80
 to_port: 80
 cidr_ip: 0.0.0.0/0
 rules_egress:
 - proto: all
 cidr_ip: 0.0.0.0/0
 register: basic_firewall
- name: Launch the new EC2 Instance
local_action: ec2
 group={{ security_group }}
 instance_type={{ instance_type }}
 image={{ image }}
 wait=true
 region={{ region }}
 keypair={{ keypair }}
 count={{ count }}
register: ec2
- name: Add Tagging to EC2 instance
local_action:
 ec2_tag resource={{ item.id }}
 region={{ region }}
 state=present
with_items: "{{ ec2.instances }}"
args:
 tags:
 Name: k8cluster
- name: associate new elastic IPs with each of the instances
ec2_eip:
 device_id: "{{ item }}"
 region: "{{ region }}"
with_items: "{{ ec2.instance_ids }}"
- name: setup a simple load balancer
ec2_elb_lb:
 name: myelb
 state: present

```

region: "{{ region }}"
zones:
  - "{{ zones }}"
listeners:
  - protocol: http
    load_balancer_port: 80
    instance_port: 80
register: myelb

- name: add the webserver to the load balancer
  local_action: ec2_elb
  args:
    instance_id: "{{ item }}"
    ec2_elbs: myelb
    state: present
    region: "{{ region }}"
  with_items: "{{ ec2.instance_ids }}"

```

Change content **Highlighted in yellow** accordingly.

1.c.3 Execute the playbook

```
$ sudo ansible-playbook create_ec2.yml
```

Expected output:

```

TASK [Create security group] *****
/usr/lib/python2.7/dist-packages/requests/__init__.py:80: RequestsDependencyWarning: urllib3 (1.26.7) or chardet (3.0.4) doesn't match a supported version!
  RequestsDependencyWarning)
[DEPRECATION WARNING]: Distribution Ubuntu 18.04 on host localhost should use /usr/bin/python3, but is using /usr/bin/python for backward compatibility with prior Ansible releases. A future Ansible release will default to using the discovered platform python for this host. See https://docs.ansible.com/ansible/2.9/reference_appendices/interpreter_discovery.html for more information. This feature will be removed in version 2.12. Deprecation warnings can be disabled by setting deprecation_warnings=False in ansible.cfg.
ok: [localhost -> localhost]

TASK [Launch the new EC2 Instance] *****
/usr/lib/python2.7/dist-packages/requests/__init__.py:80: RequestsDependencyWarning: urllib3 (1.26.7) or chardet (3.0.4) doesn't match a supported version!
  RequestsDependencyWarning)
changed: [localhost -> localhost]
/usr/lib/python2.7/dist-packages/requests/__init__.py:80: RequestsDependencyWarning: urllib3 (1.26.7) or chardet (3.0.4) doesn't match a supported version!
  RequestsDependencyWarning)

TASK [Add Tagging to EC2 instance] *****
changed: [localhost -> localhost] => (item={u'ramdisk': None, u'kernel': None, u'root_device_type': u'ebs', u'private_dns_name': u'ip-172-31-30-103.ap-southeast-1.compute.internal', u'block_device_mapping': {u'/dev/sda1': {u'status': u'attached', u'delete_on_termination': True, u'volume_id': u'vol-0dd8ede9bac0a5ffc'}}, u'key_name': u'06Dec', u'public_ip': u'18.141.164.125', u'image_id': u'ami-0907c2c44ea451f84', u'tenancy': u'default', u'private_ip': u'172.31.30.103', u'groups': {u'sg-0496dfb419b1e463c': u'k8cluster-security-group'}, u'public_dns_name': u'ec2-18-141-164-125.ap-southeast-1.compute.amazonaws.com', u'state_code': 16, u'id': u'i-08f4444cd41a430ec', u'tags': {u'': {u'ap-southeast-1b', u'ami_launch_index': u'1', u'dns_name': u'ec2-18-141-164-125.ap-southeast-1.compute.amazonaws.com', u'region': u'ap-southeast-1', u'ebs_optimized': False, u'launch_time': u'2021-12-06T15:15:32.000Z', u'instance_type': u't2.micro', u'state': u'running', u'root_device_name': u'/dev/sda1', u'hypervisor': u'xen', u'virtualization_type': u'hvm', u'architecture': u'x86_64'}}})
changed: [localhost -> localhost] => (item={u'ramdisk': None, u'kernel': None, u'root_device_type': u'ebs', u'private_dns_name': u'ip-172-31-21-6.ap-southeast-1.compute.internal', u'block_device_mapping': {u'/dev/sda1': {u'status': u'attached', u'delete_on_termination': True, u'volume_id': u'vol-058e0869759a04b05'}}, u'key_name': u'06Dec', u'public_ip': u'3.1.5.191', u'image_id': u'ami-0907c2c44ea451f84', u'tenancy': u'default', u'private_ip': u'172.31.21.6', u'groups': {u'sg-0496dfb419b1e463c': u'k8cluster-security-group'}, u'public_dns_name': u'ec2-3-1-5-191.ap-southeast-1.compute.amazonaws.com', u'state_code': 16, u'id': u'i-0c1e230a298f093', u'tags': {u'': {u'ap-southeast-1b', u'ami_launch_index': u'0', u'dns_name': u'ec2-3-1-5-191.ap-southeast-1.compute.amazonaws.com', u'region': u'ap-southeast-1', u'ebs_optimized': False, u'launch_time': u'2021-12-06T15:15:32.000Z', u'instance_type': u't2.micro', u'state': u'running', u'root_device_name': u'/dev/sda1', u'hypervisor': u'xen', u'virtualization_type': u'hvm', u'architecture': u'x86_64'}}})

TASK [associate new elastic IPs with each of the instances] *****
changed: [localhost] => (item=i-08f4444cd41a430ec)
changed: [localhost] => (item=i-0c1e230a298f093)

TASK [setup a simple load balancer] *****
ok: [localhost]

```

1.c.4 On the AWS Console, search for the ec2 service.

Expected output:

k8cluster	i-08f4444cd41a430ec	 Running		t2.micro	 2/2 checks passed	No alarms	+	ap-southeast-1b	ec2-52-220-112-84.ap-...	52.220.112.84	52.220.112.84	-	disabled	k8cluster-security-group
k8cluster	i-0c1e230a298f093	 Running		t2.micro	 Initializing	No alarms	+	ap-southeast-1b	ec2-13-214-94-31.ap-s...	13.214.94.31	13.214.94.31	-	disabled	k8cluster-security-group

Task 2 – Setting up a Kubernetes Cluster

Due to the cost involved in the EC2 instances on AWS, the Kubernetes Cluster is setup using 3 ubuntu machine on VirtualBox.


The master node: 192.168.56.51

The worker1 node: 192.168.56.52

The worker2 node: 192.168.56.53

Each node is connected to the NAT network adapter for internet access and Host-Only network adaptor for communication between the Host Machine and among the three nodes.

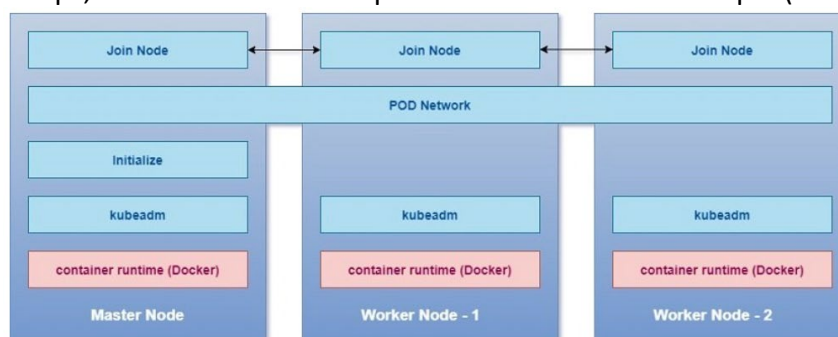
In general, we will have the following setup for the master node and all the worker nodes.



General Steps

- 1) Have multiple nodes to be formed as a kubernetes cluster
- 2) Install a container runtime on all nodes, here we are using Docker
- 3) Install kubeadm on all the nodes
- 4) Initialize the master server
- 5) Create a POD Network (or cluster network) for communication between the master node and all worker node
- 6) For all worker nodes, join node to the master node

Following the steps, we will have the setup as shown below. From step 1 (bottom) to step 6 (Top).



Pre-requisites	
Step	Description
1.a	<u>Install a Container Runtime – Docker CE</u> On all the nodes, we need to install a container runtime such as Docker. ** Repeat this for each nodes <u>Installing docker</u> <pre>\$ sudo apt-get update \$ sudo apt-get install -y docker.ce</pre>

	<p>To verify the installation:</p> <pre>\$ docker --version</pre>																					
1.b	<p><u>Configuring Docker</u></p> <p>Create a new file <code>/etc/docker/daemon.json</code> and add the following content in the file.</p> <pre>\$ sudo vi /etc/docker/daemon.json</pre> <pre>{ "exec-opts": ["native.cgroupdriver=systemd"] }</pre> <p><i>daemon.json</i></p> <p>**On the master node: Setup a local registry</p> <pre>docker run -d -p 5000:5000 --restart=always --name registry registry:2</pre> <pre>\$ sudo nano /etc/docker/daemon.json</pre> <pre>{ "exec-opts": ["native.cgroupdriver=systemd"], "insecure-registries":["192.168.56.51:5000"] }</pre> <p><i>daemon.json</i></p>																					
1.c	<p><u>Restart Docker Service</u></p> <pre>\$ service docker restart \$ service docker status</pre>																					
2.a	<p><u>Disable Swap</u></p> <p>It is required to disable swap memory for <code>kubelet</code> to work properly.</p> <p>** Repeat this for each nodes</p> <pre>\$ sudo swapoff -a</pre> <pre>ubuntu-master@ubuntumaster:~\$ free -m</pre> <table><thead><tr><th></th><th>total</th><th>used</th><th>free</th><th>shared</th><th>buff/cache</th><th>available</th></tr></thead><tbody><tr><td>Mem:</td><td>1979</td><td>519</td><td>99</td><td>7</td><td>1360</td><td>1312</td></tr><tr><td>Swap:</td><td>0</td><td>0</td><td>0</td><td></td><td></td><td></td></tr></tbody></table> <p>Output:</p>		total	used	free	shared	buff/cache	available	Mem:	1979	519	99	7	1360	1312	Swap:	0	0	0			
	total	used	free	shared	buff/cache	available																
Mem:	1979	519	99	7	1360	1312																
Swap:	0	0	0																			
2.b	<p>To ensure that swap is not re-assigned afer node reboot, we need to comment out the swap filesystem entry from <code>/etc/fstab</code></p> <pre>#/dev/mapper/rhel-swap swap swap defaults 0 0</pre>																					
3.a	<p><u>Disable SELinux</u></p> <p>We need to disable selinux or change it for Permissive mode to allow containers to access the host filesystem which is needed by pod networks.</p> <p>** Repeat this for each nodes</p> <pre>\$ setenforce 0</pre> <p>To verify:</p> <pre>\$ getenforce</pre> <p>Output: <code>Permissive</code></p>																					

3.b	<p>To make the changes persistent after reboot:</p> <pre>\$ sed -i --follow-symlinks 's/SELINUX=enforcing/SELINUX=disabled/g' /etc/sysconfig/selinux</pre>																										
4.a	<p><u>Enable Firewall</u></p> <p>We also need to enable certain pre-defined ports on the master and worker nodes.</p> <p>** On the Master Node</p> <p>Requires all workers nodes to access the ports</p> <table border="1" data-bbox="204 400 1326 763"> <thead> <tr> <th>Ports</th><th>Purpose</th></tr> </thead> <tbody> <tr> <td>6443/tcp</td><td>For Kubernetes API</td></tr> <tr> <td>6783/tcp</td><td>For weave-net</td></tr> <tr> <td>6783/udp</td><td>For weave-net</td></tr> <tr> <td>6784/udp</td><td>For weave-net</td></tr> <tr> <td>2379/tcp – 2380/tp</td><td>For etcd server client API</td></tr> <tr> <td>10250/tcp</td><td>For Kubelet API</td></tr> <tr> <td>10251/tcp</td><td>For kube-scheduler</td></tr> <tr> <td>10252tcp</td><td>For kube-controller manager</td></tr> </tbody> </table> <p>** On all Worker Nodes</p> <p>Requires all other nodes to access the ports</p> <table border="1" data-bbox="204 882 1326 1039"> <thead> <tr> <th>Ports</th><th>Purpose</th></tr> </thead> <tbody> <tr> <td>6783/tcp</td><td>For weave-net</td></tr> <tr> <td>6783/udp</td><td>For weave-net</td></tr> <tr> <td>6784/udp</td><td>For weave-net</td></tr> </tbody> </table>	Ports	Purpose	6443/tcp	For Kubernetes API	6783/tcp	For weave-net	6783/udp	For weave-net	6784/udp	For weave-net	2379/tcp – 2380/tp	For etcd server client API	10250/tcp	For Kubelet API	10251/tcp	For kube-scheduler	10252tcp	For kube-controller manager	Ports	Purpose	6783/tcp	For weave-net	6783/udp	For weave-net	6784/udp	For weave-net
Ports	Purpose																										
6443/tcp	For Kubernetes API																										
6783/tcp	For weave-net																										
6783/udp	For weave-net																										
6784/udp	For weave-net																										
2379/tcp – 2380/tp	For etcd server client API																										
10250/tcp	For Kubelet API																										
10251/tcp	For kube-scheduler																										
10252tcp	For kube-controller manager																										
Ports	Purpose																										
6783/tcp	For weave-net																										
6783/udp	For weave-net																										
6784/udp	For weave-net																										
Install Kubernetes Component																											
<p>You will install these packages on all of your machines:</p> <ul style="list-style-type: none"> kubeadm: the command to bootstrap the cluster. kubelet: the component that runs on all of the machines in your cluster and does things like starting pods and containers. kubectl: the command line util to talk to your cluster. 																											
On the Master Node																											
Step	Description																										
2.a.1	<p style="text-align: center;">Install Kubernetes Component</p> <pre>\$ sudo apt-get update</pre> <pre>\$ sudo apt-get install -y apt-transport-https ca-certificates curl</pre> <pre>\$ sudo curl -fsSLo /usr/share/keyrings/kubernetes-archive-keyring.gpg https://packages.cloud.google.com/apt/doc/apt-key.gpg</pre> <pre>\$ echo "deb [signed-by=/usr/share/keyrings/kubernetes-archive-keyring.gpg] https://apt.kubernetes.io/ kubernetes-xenial main" sudo tee /etc/apt/sources.list.d/kubernetes.list</pre> <pre>\$ sudo apt-get update</pre> <pre>\$ sudo apt-get install -y kubelet kubeadm kubectl</pre> <p><u>To verify:</u></p> <pre>\$ kubeadm version -o short</pre> <p>Output:</p>																										

```
ubuntu-master@ubuntumaster:~$ kubeadm version -o short
v1.22.3
```

```
$ kubectl version --short --client
```

Output:

```
ubuntu-master@ubuntumaster:~$ kubectl version --short --client
Client Version: version.Info{Major:"1", Minor:"22", GitVersion:"v1.22.3", GitCommit:"c92036820499fedefec0f847e2054d824aea6cd1", GitTreeState:"clean", BuildDate:"2021-10-27T18:41:28Z", GoVersion:"go1.16.9", Compiler:"gc", Platform:"linux/amd64"}
```

2.a.2

Initialize Master Node

```
$ kubeadm init --apiserver-advertise-address 192.168.56.51 --pod-network-cidr=10.244.0.0/16
```

Output:

```
Your Kubernetes control-plane has initialized successfully!

To start using your cluster, you need to run the following as a regular user:

  mkdir -p $HOME/.kube
  sudo cp -i /etc/kubernetes/admin.conf $HOME/.kube/config
  sudo chown $(id -u):$(id -g) $HOME/.kube/config

Alternatively, if you are the root user, you can run:

  export KUBECONFIG=/etc/kubernetes/admin.conf

You should now deploy a pod network to the cluster.
Run "kubectl apply -f [podnetwork].yaml" with one of the options listed at:
https://kubernetes.io/docs/concepts/cluster-administration/addons/

Then you can join any number of worker nodes by running the following on each as root:
```

```
kubeadm join 10.0.2.4:6443 --token 4wkd19.res0ndk41x38nc5x \
--discovery-token-ca-cert-hash sha256:1a61c915a43d77a17ccaebd00f548a557e06efb09aedaf32d000ed3d121f1e
```

Then execute the following commands:

```
$ mkdir -p $HOME/.kube
$ sudo cp -i /etc/kubernetes/admin.conf $HOME/.kube/config
$ sudo chown $(id -u):$(id -g) $HOME/.kube/config
```

```
ubuntu-master@ubuntumaster:~$ kubectl get nodes
NAME                STATUS    ROLES                  AGE     VERSION
ubuntumaster        Ready     control-plane,master   6m9s    v1.22.3
```

Note: You should now see a single node **master** deployed

- Copy the **kubeadm join command** that you can see on the screen of your master node after running **kubeadm init** command
- Run the **kubeadm join** command on every worker nodes as a root user to form the kubernetes cluster

2.a.3

Apply the weave-net

```
$ kubectl apply -f "https://cloud.weave.works/k8s/net?k8s-version=$(kubectl version | base64 | tr -d '\n')"
```

Output:

```
ubuntu-master@ubuntumaster:~$ kubectl apply -f "https://cloud.weave.works/k8s/net?k8s-version=$(kubectl version | base64 | tr -d '\n')"
```

```
serviceaccount/weave-net created
clusterrole.rbac.authorization.k8s.io/weave-net created
clusterrolebinding.rbac.authorization.k8s.io/weave-net created
role.rbac.authorization.k8s.io/weave-net created
rolebinding.rbac.authorization.k8s.io/weave-net created
daemonset.apps/weave-net created
```

2.a.4

Verify the setup

```
$ kubectl get nodes -o wide
```

Output:

```
ubuntu-master@ubuntu-master:~$ kubectl get nodes -o wide
NAME          STATUS    ROLES          AGE      VERSION   INTERNAL-IP   EXTERNAL-IP   OS-IMAGE             KERNEL-VERSION   CONTAINER-RUNTIME
ubuntu-master Ready    control-plane,master 3m49s    v1.22.3    10.0.2.4      <none>         Ubuntu 20.04.3 LTS   5.11.0-40-generic docker://20.10.7
```

On all Worker Nodes**Step****Description**

2.b.1

Install Kubernetes Component

```
$ sudo apt-get update

$ sudo apt-get install -y apt-transport-https ca-certificates curl

$ sudo curl -fsSLo /usr/share/keyrings/kubernetes-archive-keyring.gpg
https://packages.cloud.google.com/apt/doc/apt-key.gpg

$ echo "deb [signed-by=/usr/share/keyrings/kubernetes-archive-keyring.gpg]
https://apt.kubernetes.io/ kubernetes-xenial main" | sudo tee
/etc/apt/sources.list.d/kubernetes.list

$ sudo apt-get update

$ sudo apt-get install -y kubelet kubeadm kubectl
```

2.b.2

Join Worker Node to Master node

```
$ sudo kubeadm join 10.0.2.4:6443 --token 4wkd19.res0ndk41x38nc5x \
--discovery-token-ca-cert-hash
sha256:1a61c915a43d77a17ccaebd00fbe548a557e06efb09aedaf32d000ed3d121f1e
```

Note: Execute the kubeadm join command obtain from your screen in step 5b

Output:

```
ubuntu-worker1@ubuntu-worker1:~$ sudo kubeadm join 10.0.2.4:6443 --token 4wkd19.res0ndk41x38nc5x --discovery-token-ca-cert-hash sha256:1a61c915a43d77a17ccaebd00fbe548a557e06efb09aedaf32d000ed3d121f1e
[preflight] Running pre-flight checks
[preflight] Reading configuration from the cluster...
[preflight] FYI: You can look at this config file with 'kubectl -n kube-system get cm kubeadm-config -o yaml'
[kubelet-start] Writing kubelet configuration to file "/var/lib/kubelet/config.yaml"
[kubelet-start] Writing kubelet environment file with flags to file "/var/lib/kubelet/kubeadm-flags.env"
[kubelet-start] Starting the kubelet
[kubelet-start] Waiting for the kubelet to perform the TLS Bootstrap...

This node has joined the cluster:
* Certificate signing request was sent to apiserver and a response was received.
* The Kubelet was informed of the new secure connection details.

Run 'kubectl get nodes' on the control-plane to see this node join the cluster.
```

2.b.3

To verify setup

On the **master** node:

```
$ kubectl get nodes -o wide
```

Output:

```
ubuntu-master@ubuntu-master:~$ kubectl get nodes -o wide
NAME          STATUS    ROLES          AGE      VERSION   INTERNAL-IP   EXTERNAL-IP   OS-IMAGE             KERNEL-VERSION   CONTAINER-RUNTIME
ubuntu-master Ready    control-plane,master 74s    v1.22.3    10.0.2.101    <none>         Ubuntu 20.04.3 LTS   5.11.0-43-generic docker://20.10.7
ubuntu-worker1 Ready    <none>         27s    v1.22.3    10.0.2.102    <none>         Ubuntu 20.04.3 LTS   5.11.0-43-generic docker://20.10.7
ubuntu-worker2 Ready    <none>         22s    v1.22.3    10.0.2.103    <none>         Ubuntu 20.04.3 LTS   5.11.0-41-generic docker://20.10.7
```

Note that now the worker node has been added. Repeat this step for all other worker nodes

Task 3 – Configuration of Kubernetes Clusters for Web Application

Setup Load Balancer

Setup MetalLB Load Balancing for Bare Metal Kubernetes

****Skip this section if the Kubernetes is deployed on any cloud service provider with their own load balancer. MetalLB is one of the load balancer solution for our kubernetes cluster provision on-prem.**

Reference: <https://metallb.universe.tf/installation/>

3.a.1 Installation by Manifest

To install MetalLB by applying the manifest

```
$ kubectl apply -f
https://raw.githubusercontent.com/metallb/metallb/v0.11.0/manifests/namespace.yaml

$ kubectl apply -f
https://raw.githubusercontent.com/metallb/metallb/v0.11.0/manifests/metallb.yaml
```

3.a.2 To deploy metalLB using a ConfigMap by the Layer2 configuration

```
apiVersion: v1
kind: ConfigMap
metadata:
  namespace: metallb-system
  name: config
data:
  config: |
    address-pools:
    - name: default
      protocol: layer2
      addresses:
      - 192.168.56.0-192.168.56.50
```

loadBalance.yaml

The ip address must be a range of IP addresses that is from the same network with the kubernetes cluster. We need to serve this range of IP addresss for load balancing.

3.a.3 Apply the configMap

```
$ kubectl apply -f loadBalance.yaml
```

Setup MongoDB Database

Database - Deploy MongoDB

3.b.1 Create Storage Class

```
apiVersion: storage.k8s.io/v1
kind: StorageClass
metadata:
  name: standard
provisioner: local
parameters:
  type: standard
reclaimPolicy: Retain
allowVolumeExpansion: true
```

storage_class.yaml

```
$ kubectl apply -f storage_class.yaml
```

To verify:

```
$ kubectl get sc
```

Expected Output:

```
ubuntu-master@ubuntu-master:~$ kubectl get sc
NAME          PROVISIONER      RECLAIMPOLICY   VOLUMEBINDINGMODE   ALLOWVOLUMEEXPANSION   AGE
standard      local            Retain          Immediate            true                   62m
```

3.b.2 Create a persistence volume

```
apiVersion: v1
kind: PersistentVolume
metadata:
  name: task-pv-volume
  labels:
    type: local
spec:
  storageClassName: standard
  capacity:
    storage: 50Gi
  accessModes:
    - ReadWriteOnce
  hostPath:
    path: "/mnt/data"
```

persistence_vol.yaml

```
$ kubectl apply -f persistence_vol.yaml
```

To verify:

```
$ kubectl get pv
```

```
ubuntu-master@ubuntu-master:~$ kubectl get pv
NAME             CAPACITY   ACCESS MODES   RECLAIM POLICY   STATUS   CLAIM             STORAGECLASS   REASON   AGE
task-pv-volume   50Gi       RWX            Retain           Bound    default/pvc       standard                          64m
```

3.b.3 Create a MongoDB service

Create the mongodb-service.yaml below:

```
apiVersion: v1
kind: Service
metadata:
  labels:
    app: mongo
  name: mongo-nodeport-svc
spec:
  ports:
    - port: 27017
      protocol: TCP
      targetPort: 27017
      nodePort: 32000
  selector:
    app: mongo
```

mongo-service.yaml

On the same directory, execute the command:

```
$ kubectl apply -f mongodb-service.yaml
```

To verify:

```
$ kubectl get services
```

```
ubuntu-master@ubuntumaster:~$ kubectl get services
NAME                TYPE          CLUSTER-IP    EXTERNAL-IP    PORT(S)          AGE
kubernetes           ClusterIP     10.96.0.1     <none>         443/TCP          71m
mongo-nodeport-svc   LoadBalancer 10.97.135.205 192.168.56.10 27017:32000/TCP  56m
```

With the load balance in place, the mongoDB can be access from other remote machines on the EXTERNAL_IP **192.168.56.10**

3.b.4 Create a persistent volume claim for mongoDB

```
apiVersion: v1
kind: PersistentVolumeClaim
metadata:
  name: pvc
spec:
  storageClassName: "standard"
  accessModes:
    - ReadWriteOnce
  volumeName: task-pv-volume
  resources:
    requests:
      storage: 1Gi
```

mongodb-pvc.yaml

To apply:

```
$ kubectl apply -f mongodb-pvc.yaml
```

To verify:

```
$ kubectl get pvc
```

```
ubuntu-master@ubuntumaster:~$ kubectl get pvc
NAME    STATUS    VOLUME          CAPACITY   ACCESS MODES   STORAGECLASS   AGE
pvc     Bound     task-pv-volume   50Gi       RW0             standard       68m
```

3.b.5 Create the mongoDB deployment

```
apiVersion: apps/v1
kind: Deployment
metadata:
  creationTimestamp: null
  labels:
    app: mongo
  name: mongo
spec:
  replicas: 1
  selector:
    matchLabels:
      app: mongo
  strategy: {}
  template:
    metadata:
```

```

creationTimestamp: null
labels:
  app: mongo
spec:
  containers:
  - image: mongo
    name: mongod-container
    args: ["--dbpath", "/data/db"]
    livenessProbe:
      exec:
        command:
        - mongo
        - --disableImplicitSessions
        - --eval
        - "db.adminCommand('ping')"
      initialDelaySeconds: 30
      periodSeconds: 10
      timeoutSeconds: 5
      successThreshold: 1
      failureThreshold: 6
    readinessProbe:
      exec:
        command:
        - mongo
        - --disableImplicitSessions
        - --eval
        - "db.adminCommand('ping')"
      initialDelaySeconds: 30
      periodSeconds: 10
      timeoutSeconds: 5
      successThreshold: 1
      failureThreshold: 6
    env:
    - name: MONGO_INITDB_ROOT_USERNAME
      value: admin #user-defined username for root user of mongoDB
    - name: MONGO_INITDB_ROOT_PASSWORD
      value: admin123 #user-defined password root user of mongoDB
    volumeMounts:
    - name: "mongo-data-dir"
      mountPath: "/data/db"
  volumes:
  - name: "mongo-data-dir"
    persistentVolumeClaim:
      claimName: "pvc"

```

mongodb-deployment.yaml

On the same directory, execute the command:

```
$ kubectl apply -f mongodb-deployment.yaml
```

To verify:

```
$ kubectl get deployment
```

```

ubuntu-master@ubuntumaster:~$ kubectl get deployments
NAME          READY    UP-TO-DATE    AVAILABLE    AGE
mongo         1/1      1             1            71m

```

Along with the external IP address created when the mongoDB service is created, we can connect to the mongoDB using the url:

mongodb://admin:admin123@192.168.56.10:27017/?authSource=admin

In general, the URI is formatted

<code>mongodb://<_username>:<password>@<external_IP_address_of_mongoDB_service>:27017/?authSource=admin</code>
--

Setup Web Application

Both the backend and the frontend of the web application is based on NodeJS. The nodeJS server is responsible for communicating with the mongoDB for CRUD tasks, as well as rendering the html files on client machines.

Prepare Web Application docker image

3.c.1 Login to docker using CLI

```
$ sudo docker login
```

When prompt, enter your docker username and password. This will allow us to push docker image build locally to docker hub.

Expected Output:

```
ubuntu-master@ubuntumaster:~$ sudo docker login
[sudo] password for ubuntu-master:
Authenticating with existing credentials...
WARNING! Your password will be stored unencrypted in /root/.docker/config.json.
Configure a credential helper to remove this warning. See
https://docs.docker.com/engine/reference/commandline/login/#credentials-store

Login Succeeded
```

3.c.2 Dockerize NodeJS Application

The NodeJS Application has the following folder structure:

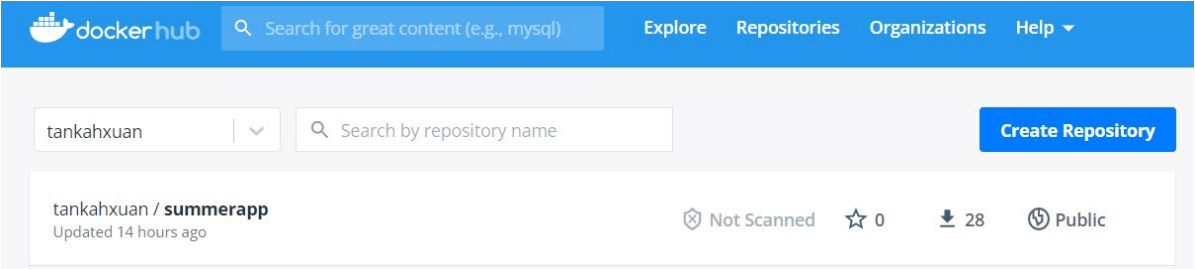
```
|— controller
|   |— app.js
|— models
|— package-lock.json
|— server.js
|— views
|   |— images
|   |— reviews.ejs
|   |— css
|   |— home.ejs
|— package.json
|— Dockerfile
```

```
FROM node:16
```

```
RUN mkdir /opt/SummerMelts
```

```
WORKDIR /opt/SummerMelts
```

```
COPY package.json ./
```


	<pre> RUN npm install COPY . . EXPOSE 3000 CMD ["node", "server.js"] </pre> <hr/> <p><i>Dockerfile</i></p> <p>The web application will be expose on port 3000.</p> <p>Build the image:</p> <pre>\$ sudo docker build -t summerapp:v1 .</pre>
3.c.3	<p><u>Re-tagging an existing local image</u></p> <pre>\$ docker tag summerapp:v1 tankahxuan/summerapp:v1</pre> <p>Note: In future, where any changes is made, execute the command to commit the changes</p> <pre>\$ docker commit <existing-container> <hub-user>/<repo-name>[:<tag>]</pre>
3.c.4	<p><u>Push the image to the docker registry</u></p> <pre>\$ docker push tankahxuan/summerapp:v1</pre> <p>On docker hub:</p>  <p>https://hub.docker.com/repository/docker/tankahxuan/summerapp</p>
Deploy Web Application on K8s Cluster	
3.c.5	<p>Create a deployment and service object on K8s cluster</p> <hr/> <pre> apiVersion: v1 kind: Service metadata: name: summerapp-backend spec: type: LoadBalancer selector: app: summerapp-app ports: - port: 3000 targetPort: 3000 protocol: 'TCP' --- apiVersion: apps/v1 kind: Deployment metadata: </pre>

```

name: summerapp-deployment
labels:
  app: summerapp
spec:
  replicas: 1
  selector:
    matchLabels:
      app: summerapp-app
      tier: summerapp-backend
  template:
    metadata:
      labels:
        app: summerapp-app
        tier: summerapp-backend
    spec:
      containers:
      - name: summerapp-container
        image: tankahxuan/summerapp:v1
        imagePullPolicy: Always
        livenessProbe:
          httpGet:
            path: /
            port: 3000
          periodSeconds: 20
          initialDelaySeconds: 10
        ports:
        - containerPort: 3000
        env:
        - name: DATABASE_URI
          value: 'mongodb://admin:admin123@192.168.56.0:27017/?authSource=admin'

```

myappDeploy.yaml

3.c.5

On the same directory, execute the command:

```
$ kubectl apply -f myappDeploy.yaml
```

To verify:

```

ubuntu-master@ubuntumaster:~/learnK8/myapp$ kubectl get deployment
NAME                READY   UP-TO-DATE   AVAILABLE   AGE
mongo                1/1     1             1           103m
summerapp-deployment 0/1     1             0           93m

```

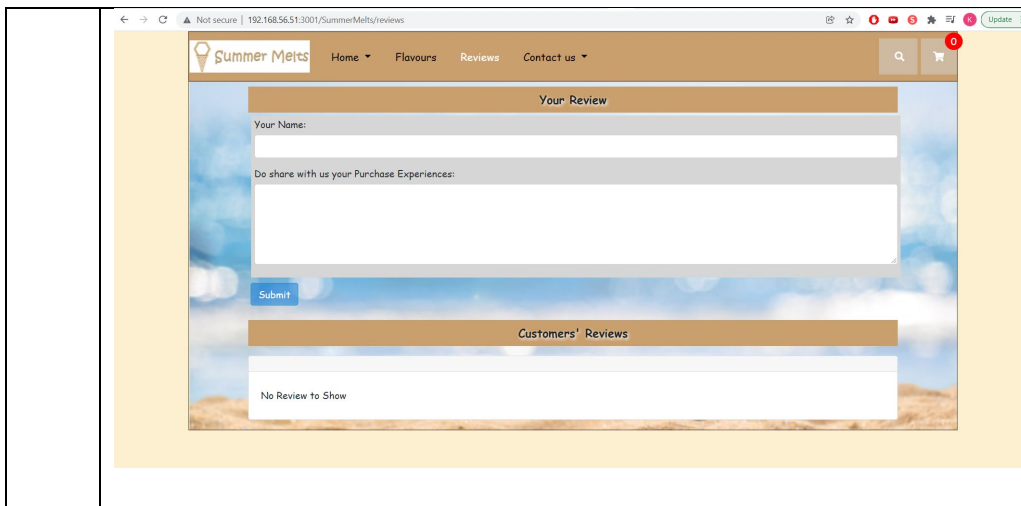
```

ubuntu-master@ubuntumaster:~/learnK8/myapp$ kubectl get services
NAME                TYPE        CLUSTER-IP      EXTERNAL-IP      PORT(S)          AGE
kubernetes           ClusterIP   10.96.0.1        <none>            443/TCP          110m
mongo-nodeport-svc   LoadBalancer 10.97.135.205    192.168.56.10    27017:32000/TCP  94m
summerapp-backend    LoadBalancer 10.98.144.68     192.168.56.11    3000:30150/TCP   93m

```

To access the web application, open a browser and enter:

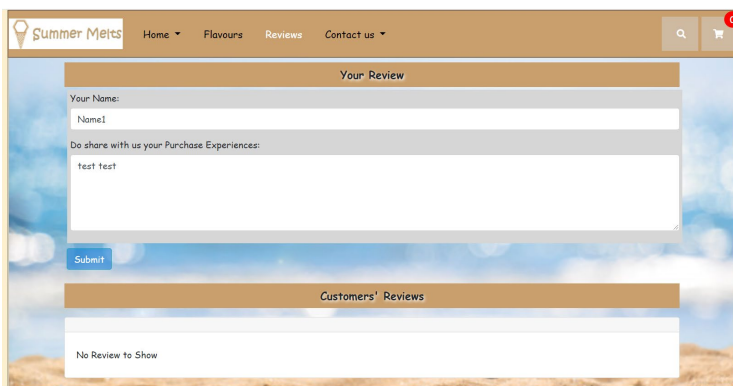
<http://192.168.56.11:3000/SummerMelts/reviews>



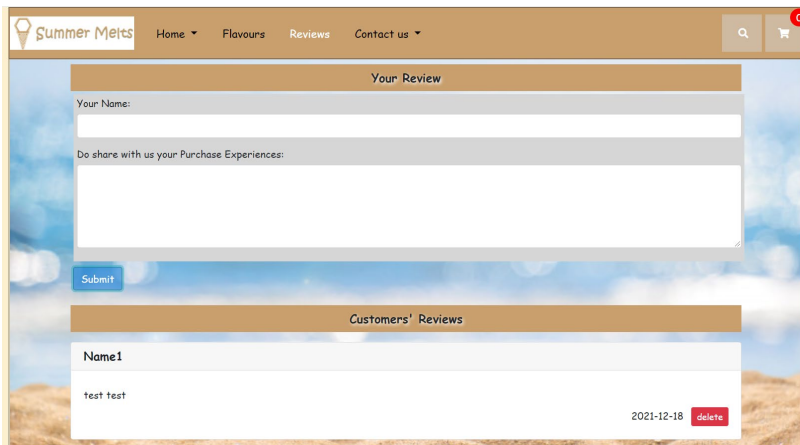
Web Application Functions

This is a simple web application for users to enter their reviews. The reviews will then be saved in mongoDB. Thereafter, all reviews in mongoDB will be retrieved and display on the same web page.

On the web page, enter any random name and describe and click the submit button.



The data will then be saved in the database and retrieve by the web application to be display under the customer's review.



In the database:

MongoDB Compass - 192.168.56.10:27017/test.reviews

ConnectViewCollectionHelp

Local

4 DBS15 COLLECTIONS

☆ FAVORITE

HOST
192.168.56.10:27017

CLUSTER
Standalone

EDITION
MongoDB 5.0.5 Community

Filter your data

> admin

> config

> local

> test

reviews

test.reviews

Documents

test.reviews Documents

DOCUMENTS 2

TOTAL SIZE 259B

AVG. SIZE 130B

INDEXES 1

TOTAL SIZE 36.9KB

AVG. SIZE 36.9KB

Documents

Aggregations

Schema

Explain Plan

Indexes

Validation

FILTER { field: 'value' }

OPTIONS

FIND

RESET

ADD DATA

VIEW

Displaying documents 1 - 1 of 1

REFRESH

_id: ObjectId("61bd8d91badaf01f62ea2dd4")

review_id: "Name1_2021-12-18"

name: "Name1"

review_text: "test test"

review_date: "2021-12-18"

__v: 0

Task 4 – Create Users in the Kubernetes Cluster

Create Client Certificate

4.a.1 Create a client key

```
$ openssl genrsa -out ubuntu-master.key 2048
```

```
ubuntu-master@ubuntumaster:~$ openssl genrsa -out ubuntu-master.key 2048
Generating RSA private key, 2048 bit long modulus (2 primes)
.....+++++
.....+++++
e is 65537 (0x010001)
```

4.a.2 Create a certificate signing request

```
$ openssl req -new -key ubuntu-master.key -out ubuntu-master.csr
```

4.a.3 Copy the CA certification key to the root folder

Navigate to the root folder: /home/ubuntu-master/

```
$ cd /etc/kubernetes/pki
```

```
$ sudo cp ca.crt ca.key /home/ubuntu-master/
```

```
ubuntu-master@ubuntumaster:~$ ls
ca.crt  ca.srl  Documents  learnK8  Pictures  rolebinding.yaml  Templates  ubuntu-master.csr  Videos
ca.key  Desktop  Downloads  Music    Public    role.yaml         ubuntu-master.crt  ubuntu-master.key
```

4.a.4 Sign the user key and request with cluster certificate

```
$ sudo openssl x509 -req -in ubuntu-master.csr -CA ca.crt -CAkey ca.key -CAcreateserial -out ubuntu-master.crt -days 300
```

Add user credentials to the kubeconfig file

4.b.1 Add user credentials

```
$ kubectl config set-credentials ubuntu-master --client-certificate=ubuntu-master.crt --client-key=ubuntu-master.key
```

Create Role and Role Binding

4.c.1 Create a role with permission to create, list, update and delete pods

```
kind: Role
apiVersion: rbac.authorization.k8s.io/v1
metadata:
  name: get-pods
rules:
- apiGroups: ["*"]
  resources: ["pods"]
  verbs: ["list", "get", "update", "delete"]
```

role.yaml

```
$ kubectl apply -f role.yaml
```

4.c.2 Bind the role to a user

```
$ kubectl apply -f rolebinding.yaml
```

NAME	READY	STATUS	RESTARTS	AGE
coredns-78fcd69978-bsksb	1/1	Running	0	69m
coredns-78fcd69978-nr272	1/1	Running	0	69m
etcd-ubuntu-master	1/1	Running	6	70m
kube-apiserver-ubuntu-master	1/1	Running	2	70m
kube-controller-manager-ubuntu-master	1/1	Running	4	70m
kube-proxy-227sv	1/1	Running	0	69m
kube-proxy-7flhr	1/1	Running	0	69m
kube-proxy-sprx6	1/1	Running	0	69m
kube-scheduler-ubuntu-master	1/1	Running	4	70m
weave-net-fwmw8	2/2	Running	0	69m
weave-net-jjk22	2/2	Running	0	69m
weave-net-nz86b	2/2	Running	0	69m

5.a.3	Obtain the advertise-client-url of the etcd-pod
	<pre>\$ export advertise_url=https://192.168.56.51:2379 \$ echo \$advertise_url</pre>
	<pre>ubuntu-master@ubuntumaster:~\$ kubectl describe pod etcd-ubuntumaster -n kube-system Name: etcd-ubuntumaster Namespace: kube-system Priority: 2000001000 Priority Class Name: system-node-critical Node: ubuntu-master/10.0.2.101 Start Time: Sat, 18 Dec 2021 21:35:38 +0800 Labels: component=etcd tier=control-plane Annotations: kubeadm.kubernetes.io/etcd.advertise-client-urls: https://192.168.56.51:2379 kubernetes.io/config.hash: b2a8f224c1fc27870bae1f3185ef904d kubernetes.io/config.mirror: b2a8f224c1fc27870bae1f3185ef904d kubernetes.io/config.seen: 2021-12-18T21:35:38.731450760+08:00 kubernetes.io/config.source: file seccomp.security.alpha.kubernetes.io/pod: runtime/default Status: Running IP: 10.0.2.101 IPs: IP: 10.0.2.101 Controlled By: Node/ubuntumaster Containers: etcd: Container ID: docker://ba7361cc30b36ae638e6910806cb8b8fac71fc32ad3f97ab8849abd9980ae3e5 Image: k8s.gcr.io/etcd:3.5.0-0 Image ID: docker-pullable://k8s.gcr.io/etcd@sha256:9ce33ba33d8e738a5b85ed50b5080ac746dececd4a7496c550927a7a19ca3b6d Port: <none> Host Port: <none> Command: etcd --advertise-client-urls=https://192.168.56.51:2379 --cert-file=/etc/kubernetes/pki/etcd/server.crt --client-cert-auth=true --data-dir=/var/lib/etcd --initial-advertise-peer-urls=https://192.168.56.51:2380 --initial-cluster=ubuntumaster=https://192.168.56.51:2380 --key-file=/etc/kubernetes/pki/etcd/server.key --listen-client-urls=https://127.0.0.1:2379,https://192.168.56.51:2379 --listen-metrics-urls=http://127.0.0.1:2381</pre> <p>Here, the advertise-client-urls is <i>https://192.168.56.51:2379</i></p>
5.a.4	Set the advertis-client-url to environment variable
	<pre>ubuntu-master@ubuntumaster:~\$ export advertise_url=https://192.168.56.51:2379 ubuntu-master@ubuntumaster:~\$ echo \$advertise_url https://192.168.56.51:2379</pre>
5.a.5	Create the backup
	<p>Execute the following command to create a snapshot of the etcd database, named "etcd_backup.db"</p> <pre>sudo ETCDCTL_API=3 etcdctl \ --endpoints \$advertise_url \ --cacert /etc/kubernetes/pki/etcd/ca.crt \ --key /etc/kubernetes/pki/etcd/server.key \ --cert /etc/kubernetes/pki/etcd/server.crt snapshot save etcd_backup.db</pre>

Task 6 – Auto Scaling

Reset kubernetes cluster

If there is a need to reconfigure or reset your existing Kubernetes cluster, follow the steps below.

a Create the horizontal scaling object

```
apiVersion: autoscaling/v2beta2
kind: HorizontalPodAutoscaler
metadata:
  name: horizontal-scale
spec:
  scaleTargetRef:
    apiVersion: apps/v1
```

	<pre> kind: Deployment name: summerapp-deployment minReplicas: 1 maxReplicas: 10 metrics: - type: Resource resource: name: cpu target: type: Utilization averageUtilization: 50 </pre>
	<pre> autoScale.yaml </pre>
	<pre> \$ kubectl apply -f autoScale.yaml </pre>

Appendix – To reset kubernetes cluster

Reset kubernetes cluster	
If there is a need to reconfigure or reset your existing Kubernetes cluster, follow the steps below.	
a	<p>Reset kubernetes cluster using kubeadm</p> <pre> \$ sudo kubeadm reset -f </pre> <p>Output:</p> <pre> ubuntu-master@ubuntu-master:~\$ sudo kubeadm reset [reset] WARNING: Changes made to this host by 'kubeadm init' or 'kubeadm join' will be reverted. [reset] Are you sure you want to proceed? [y/N]: y [preflight] Running pre-flight checks W1124 11:25:59.094585 13080 removeetcdmember.go:80] [reset] No kubeadm config, using etcd pod spec to get data directory [reset] No etcd config found. Assuming external etcd [reset] Please, manually reset etcd to prevent further issues [reset] Stopping the kubelet service [reset] Unmounting mounted directories in "/var/lib/kubelet" [reset] Deleting contents of config directories: [/etc/kubernetes/manifests /etc/kubernetes/pki] [reset] Deleting files: [/etc/kubernetes/admin.conf /etc/kubernetes/kubelet.conf /etc/kubernetes/bootstrap-kubelet.conf /etc/kubernetes/controller-manager.conf /etc/kubernetes/scheduler.conf] [reset] Deleting contents of stateful directories: [/var/lib/kubelet /var/lib/docker/shim /var/run/kubernetes /var/lib/cni] The reset process does not clean CNI configuration. To do so, you must remove /etc/cni/net.d The reset process does not reset or clean up iptables rules or IPVS tables. If you wish to reset iptables, you must do so manually by using the "iptables" command. If your cluster was setup to utilize IPVS, run ipvsadm --clear (or similar) to reset your system's IPVS tables. The reset process does not clean your kubeconfig files and you must remove them manually. Please, check the contents of the \$HOME/.kube/config file. </pre>
b	<p>Remove all the data from all below location</p> <pre> \$ sudo rm -rf /etc/cni /etc/kubernetes /var/lib/docker/shim /var/lib/etcd /var/lib/kubelet /var/run/kubernetes ~/.kube/* </pre> <p>Restart the docker service</p> <pre> \$ systemctl restart docker </pre>

Reset kubernetes cluster

If there is a need to reconfigure or reset your existing Kubernetes cluster, follow the steps below.

a	<div data-bbox="555 103 1093 134" data-label="Section-Header"> <h3>Reset kubernetes cluster using kubeadm</h3> </div> <div data-bbox="156 165 531 197" data-label="Text"> <pre>\$ sudo kubeadm reset -f</pre> </div> <div data-bbox="156 210 261 241" data-label="Text"> <p>Output:</p> </div> <div data-bbox="156 264 1430 705" data-label="Text"> <pre>ubuntu-master@ubuntu-master:~\$ sudo kubeadm reset [reset] WARNING: Changes made to this host by 'kubeadm init' or 'kubeadm join' will be reverted. [reset] Are you sure you want to proceed? [y/N]: y [preflight] Running pre-flight checks W1124 11:25:59.094585 13080 removeetcdmember.go:80] [reset] No kubeadm config, using etcd pod spec to get data directory [reset] No etcd config found. Assuming external etcd [reset] Please, manually reset etcd to prevent further issues [reset] Stopping the kubelet service [reset] Unmounting mounted directories in "/var/lib/kubelet" [reset] Deleting contents of config directories: [/etc/kubernetes/manifests /etc/kubernetes/pki] [reset] Deleting files: [/etc/kubernetes/admin.conf /etc/kubernetes/kubelet.conf /etc/kubernetes/bootstrap-kubelet.conf /etc/kubernetes/controller-manager.conf /etc/kubernetes/scheduler.conf] [reset] Deleting contents of stateful directories: [/var/lib/kubelet /var/lib/docker/shim /var/run/kubernetes /var/lib/cni] The reset process does not clean CNI configuration. To do so, you must remove /etc/cni/net.d The reset process does not reset or clean up iptables rules or IPVS tables. If you wish to reset iptables, you must do so manually by using the "iptables" command. If your cluster was setup to utilize IPVS, run ipvsadm --clear (or similar) to reset your system's IPVS tables. The reset process does not clean your kubeconfig files and you must remove them manually. Please, check the contents of the \$HOME/.kube/config file.</pre> </div>
b	<div data-bbox="533 797 1115 828" data-label="Section-Header"> <h3>Remove all the data from all below location</h3> </div> <div data-bbox="156 878 1321 938" data-label="Text"> <pre>\$ sudo rm -rf /etc/cni /etc/kubernetes /var/lib/docker/shim /var/lib/etcd /var/lib/kubelet /var/run/kubernetes ~/.kube/*</pre> </div> <div data-bbox="651 1021 995 1052" data-label="Section-Header"> <h3>Restart the docker service</h3> </div> <div data-bbox="156 1102 577 1133" data-label="Text"> <pre>\$ systemctl restart docker</pre> </div>